DVA_lab-5_mahalakshmi18

March 28, 2021

0.0.1 Lab5. Pandas Concatenate, Merge and Join

[2]:
```
import pandas as pd            # Import necessary modules
import matplotlib.pyplot as plt
%matplotlib inline
```

First column should be used as the row index by passing the argument index_col=0

[3]:
```
north_america = pd.read_csv('north_america_2000_2010.csv', index_col=0)
south_america = pd.read_csv('south_america_2000_2010.csv', index_col=0)
```

[4]:
```
north_america        #north_america #(UNCOMMENT AND SEE OUTPUT)
```

[4]:

Country	2000	2001	2002	2003	2004	2005	2006	2007	2008	2009	2010
Canada	1779.0	1771.0	1754.0	1740.0	1760.0	1747	1745.0	1741.0	1735	1701.0	1703.0
Mexico	2311.2	2285.2	2271.2	2276.5	2270.6	2281	2280.6	2261.4	2258	2250.2	2242.4
USA	1836.0	1814.0	1810.0	1800.0	1802.0	1799	1800.0	1798.0	1792	1767.0	1778.0

[5]:
```
south_america        #south_america #(UNCOMMENT AND SEE OUTPUT)
```

[5]:

Country	2000	2001	2002	2003	2004	2005	2006	2007	2008	2009	2010
Chile	2263	2242	2250	2235	2232	2157	2165	2128	2095	2074	2069.6

Here, rows are countries, columns are years, and cell values are the average annual hours worked per employee.

0.0.2 Create line graphs for our yearly labor trends in north_america

[6]:
```
north_america.plot()
```

[6]:
```
<AxesSubplot:xlabel='Country'>
```

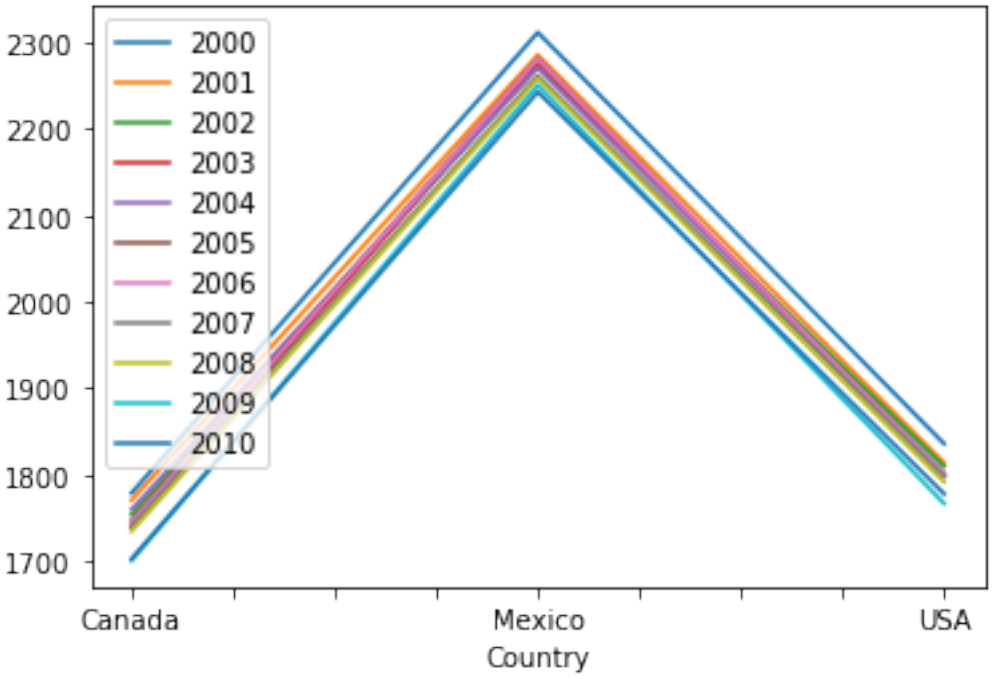

0.0.3 Plot transposed line graph of north_america dataframe, with title "Average Labor Hours Per Year"

[7]:
```
north_america.transpose().plot(title='Average Labor Hours Per Year')
plt.show()
```

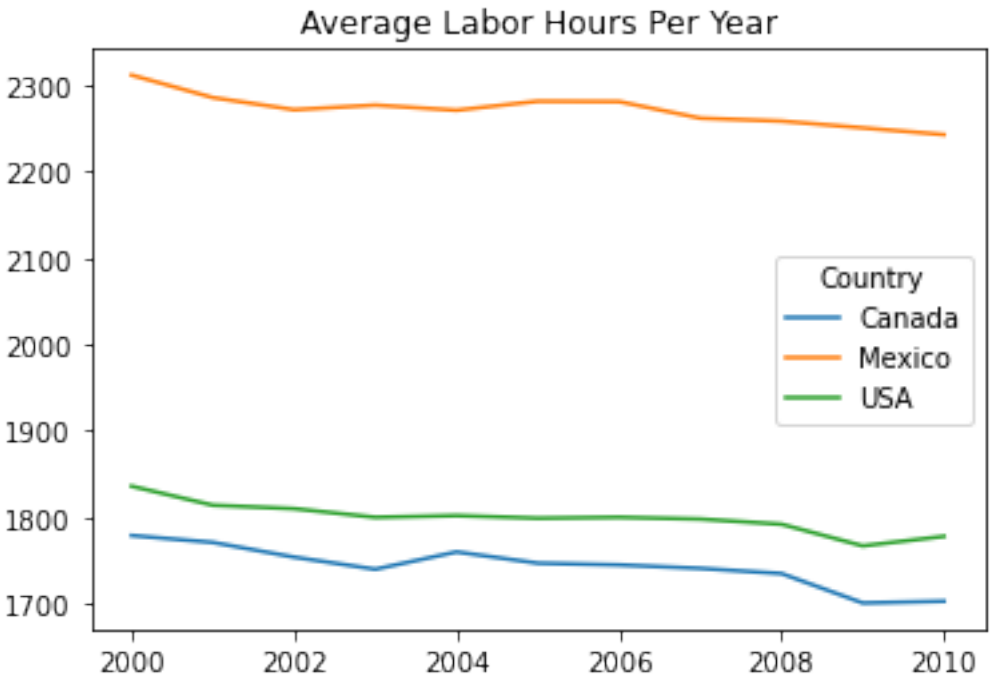

0.0.4 Similarly, plot transposed south_america dataframe with title "Average Labor Hours Per Year". Output chart is shown below

[8]:
```
south_america.transpose().plot(title='Average Labor Hours Per Year')
plt.show()
```

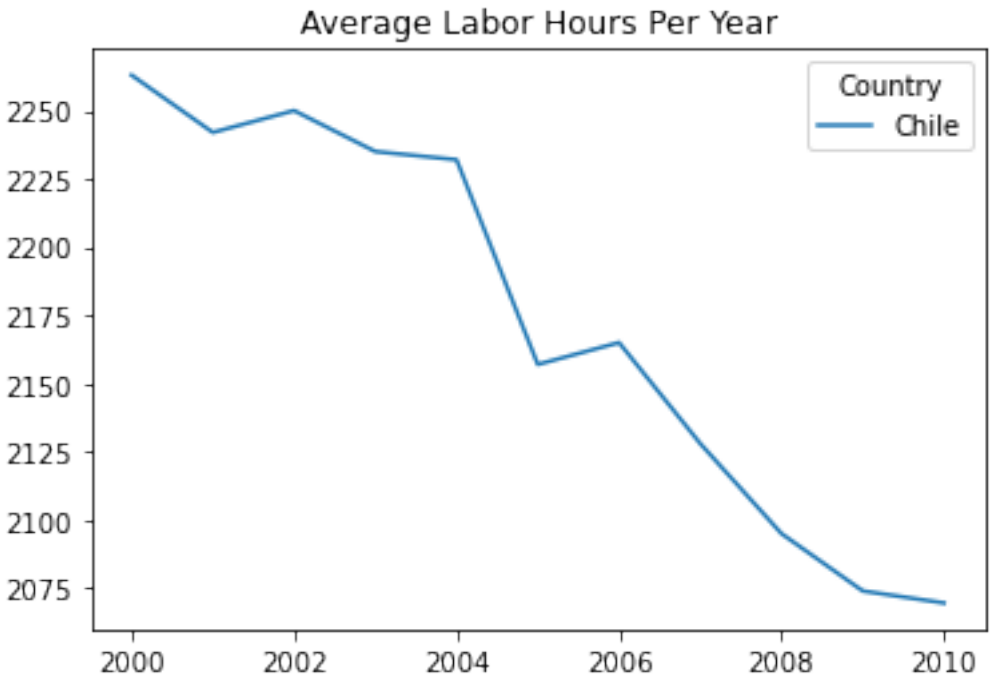

0.0.5 Concatenate America DataIt's hard to compare the average labor hours in South America versus North America.

If we were able to get all the countries into the same data frame, it would be much easier to do this camparison.

0.0.6 Concatenate north_america and south_america dataframes and store result in a dataframe,americas

[38]:
```
americas = pd.concat([north_america, south_america])
americas
```

[38]:

Country	2000	2001	2002	2003	2004	2005	2006	2007	2008	2009	2010
Canada	1779.0	1771.0	1754.0	1740.0	1760.0	1747	1745.0	1741.0	1735	1701.0	1703.0
Mexico	2311.2	2285.2	2271.2	2276.5	2270.6	2281	2280.6	2261.4	2258	2250.2	2242.4
USA	1836.0	1814.0	1810.0	1800.0	1802.0	1799	1800.0	1798.0	1792	1767.0	1778.0
Chile	2263.0	2242.0	2250.0	2235.0	2232.0	2157	2165.0	2128.0	2095	2074.0	2069.6

Now, our data collection team has sent us data files for each year from 2011 to 2015 in separate CSV files. They are americas_2011.csv , americas_2012.csv, americas_2014.csv and americas_2015.csv

0.0.7 Load the additional files

[39]:
```
americas_11 = pd.read_csv('americas_2011.csv', index_col=0)
americas_12 = pd.read_csv('americas_2012.csv', index_col=0)
americas_13 = pd.read_csv('americas_2013.csv', index_col=0)
americas_14 = pd.read_csv('americas_2014.csv', index_col=0)
americas_15 = pd.read_csv('americas_2015.csv', index_col=0)
```

[40]:
```
t=americas_11.join(americas_12)
```

[41]:
```
t=t.join(americas_13)
```

[42]:
```
t=t.join(americas_14)
```

[43]:
```
t=t.join(americas_15)
```

[44]:
```
t
```

[44]:

Country	2011	2012	2013	2014	2015
Canada	1700.0	1713.0	1707.0	1703.0	1706.0
Chile	2047.4	2024.0	2015.3	1990.1	1987.5
Mexico	2250.2	2225.8	2236.6	2228.4	2246.4
USA	1786.0	1789.0	1787.0	1789.0	1790.0

[45]:
```
americas = americas.join(t)
```

[46]:
```
americas.index.names = ['Country']
```

[47]:
```
americas
```

[47]:

Country	2000	2001	2002	2003	2004	2005	2006	2007	2008	\
Canada	1779.0	1771.0	1754.0	1740.0	1760.0	1747	1745.0	1741.0	1735	
Mexico	2311.2	2285.2	2271.2	2276.5	2270.6	2281	2280.6	2261.4	2258	
USA	1836.0	1814.0	1810.0	1800.0	1802.0	1799	1800.0	1798.0	1792	
Chile	2263.0	2242.0	2250.0	2235.0	2232.0	2157	2165.0	2128.0	2095	

Country	2009	2010	2011	2012	2013	2014	2015
Canada	1701.0	1703.0	1700.0	1713.0	1707.0	1703.0	1706.0
Mexico	2250.2	2242.4	2250.2	2225.8	2236.6	2228.4	2246.4
USA	1767.0	1778.0	1786.0	1789.0	1787.0	1789.0	1790.0
Chile	2074.0	2069.6	2047.4	2024.0	2015.3	1990.1	1987.5

0.0.8 Concatenate americas and americas_dfs dataframes and store result in americas

[48]:
```
americas_dfs = [americas]
americas = pd.concat(americas_dfs, axis=1)
```

[49]:
```
americas.index.names = ['Country']
americas
```

[49]:

Country	2000	2001	2002	2003	2004	2005	2006	2007	2008	\
Canada	1779.0	1771.0	1754.0	1740.0	1760.0	1747	1745.0	1741.0	1735	
Mexico	2311.2	2285.2	2271.2	2276.5	2270.6	2281	2280.6	2261.4	2258	
USA	1836.0	1814.0	1810.0	1800.0	1802.0	1799	1800.0	1798.0	1792	
Chile	2263.0	2242.0	2250.0	2235.0	2232.0	2157	2165.0	2128.0	2095	

Country	2009	2010	2011	2012	2013	2014	2015
Canada	1701.0	1703.0	1700.0	1713.0	1707.0	1703.0	1706.0
Mexico	2250.2	2242.4	2250.2	2225.8	2236.6	2228.4	2246.4
USA	1767.0	1778.0	1786.0	1789.0	1787.0	1789.0	1790.0
Chile	2074.0	2069.6	2047.4	2024.0	2015.3	1990.1	1987.5

0.0.9 Now, plot transposed americas dataframe

[50]:
```
americas.transpose().plot(title='Average Labor Hours Per Year')
plt.show()
```

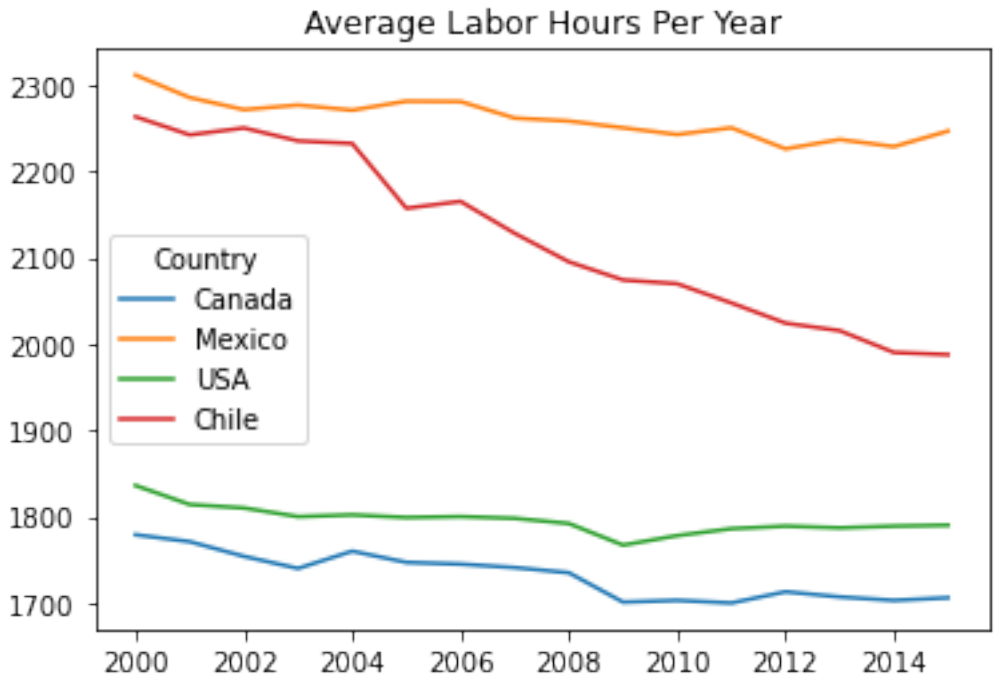

0.0.10 Appending data from other Continents

The data collection team has provided CSV files for Asia, Europe, and the South Pacific for 2000 through 2015. Let's load these files in and have a preview

[51]:
```
asia = pd.read_csv('asia_2000_2015.csv', index_col=0)
asia
```

[51]:

Country	2000	2001	2002	2003	2004	2005	2006	2007	2008	2009	2010	2011	2012	2013	2014	2015
Israel	2017	1979	1993	1974	1942	1931	1919	1931	1929	1927	1918	1920	1910	1867	1853	1858
Japan	1821	1809	1798	1799	1787	1775	1784	1785	1771	1714	1733	1728	1745	1734	1729	1719
Korea	2512	2499	2464	2424	2392	2351	2346	2306	2246	2232	2187	2090	2163	2079	2124	2113
Russia	1982	1980	1982	1993	1993	1989	1998	1999	1997	1974	1976	1979	1982	1980	1985	1978

[52]:
```
europe = pd.read_csv('europe_2000_2015.csv', index_col=0)
europe.head()
```

[52]:

Country	2000	2001	2002	2003	2004	2005	2006	2007	2008	2009	2010	2011	2012	2013	2014	2015
Austria	1807.4	1794.6	1792.2	1783.8	1786.8	1764.0	1746.2	1736.0	1728.5	1673.0	1668.6	1675.9	1652.9	1636.7	1629.4	1624.9
Belgium	1595.0	1588.0	1583.0	1578.0	1573.0	1565.0	1572.0	1577.0	1570.0	1548.0	1546.0	1560.0	1560.0	1558.0	1560.0	1541.0
Switzerland	1673.6	1635.0	1614.0	1626.8	1656.5	1651.7	1643.2	1632.7	1623.1	1614.9	1612.4	1605.4	1590.9	1572.9	1568.3	1589.7
Czech Republic	1896.0	1818.0	1816.0	1806.0	1817.0	1817.0	1799.0	1784.0	1790.0	1779.0	1800.0	1806.0	1776.0	1763.0		
Germany	1452.0	1441.9	1430.9	1424.8	1422.2	1411.3	1424.7	1424.4	1418.4	1372.7	1389.9	1392.8	1375.3	1361.7		

```
Czech Republic  1771.0  1779.0
Germany         1366.4  1371.0
```

[53]:
```
south_pacific = pd.read_csv('south_pacific_2000_2015.csv', index_col=0)
south_pacific
```

[53]:
```
                2000    2001    2002    2003    2004    2005    2006    2007  \
Country
Australia    1778.7  1736.7  1731.7  1735.8  1734.5  1729.2  1720.5  1712.5
New Zealand  1836.0  1825.0  1826.0  1823.0  1830.0  1815.0  1795.0  1774.0

                2008  2009    2010    2011    2012    2013    2014  2015
Country
Australia    1717.2  1690  1691.5  1699.5  1678.6  1662.7  1663.6  1665
New Zealand  1761.0  1740  1755.0  1746.0  1734.0  1752.0  1762.0  1757
```

If any columns were missing from the data we are trying to append, they would result in those rows having NaN values in the cells falling under the missing year columns. Let's run the append method and verify that all the countries have been sucesfully appended by printing DataFrame.index.

0.0.11 Append asia, europe and south_pacific to americas dataframe and assign to new dataframe world

[54]:
```
world = americas.append([asia, europe, south_pacific])
world.index
```

[54]:
```
Index(['Canada', 'Mexico', 'USA', 'Chile', 'Israel', 'Japan', 'Korea',
       'Russia', 'Austria', 'Belgium', 'Switzerland', 'Czech Republic',
       'Germany', 'Denmark', 'Spain', 'Estonia', 'Finland', 'France',
       'United Kingdom', 'Greece', 'Hungary', 'Ireland', 'Iceland', 'Italy',
       'Lithuania', 'Luxembourg', 'Latvia', 'Netherlands', 'Norway', 'Poland',
       'Portugal', 'Slovak Republic', 'Slovenia', 'Sweden', 'Australia',
       'New Zealand'],
      dtype='object', name='Country')
```

0.0.12 Plot, transposed world dataframe

[55]:
```
world.transpose().plot(title='Average Labor Hours Per Year')
plt.show()
```

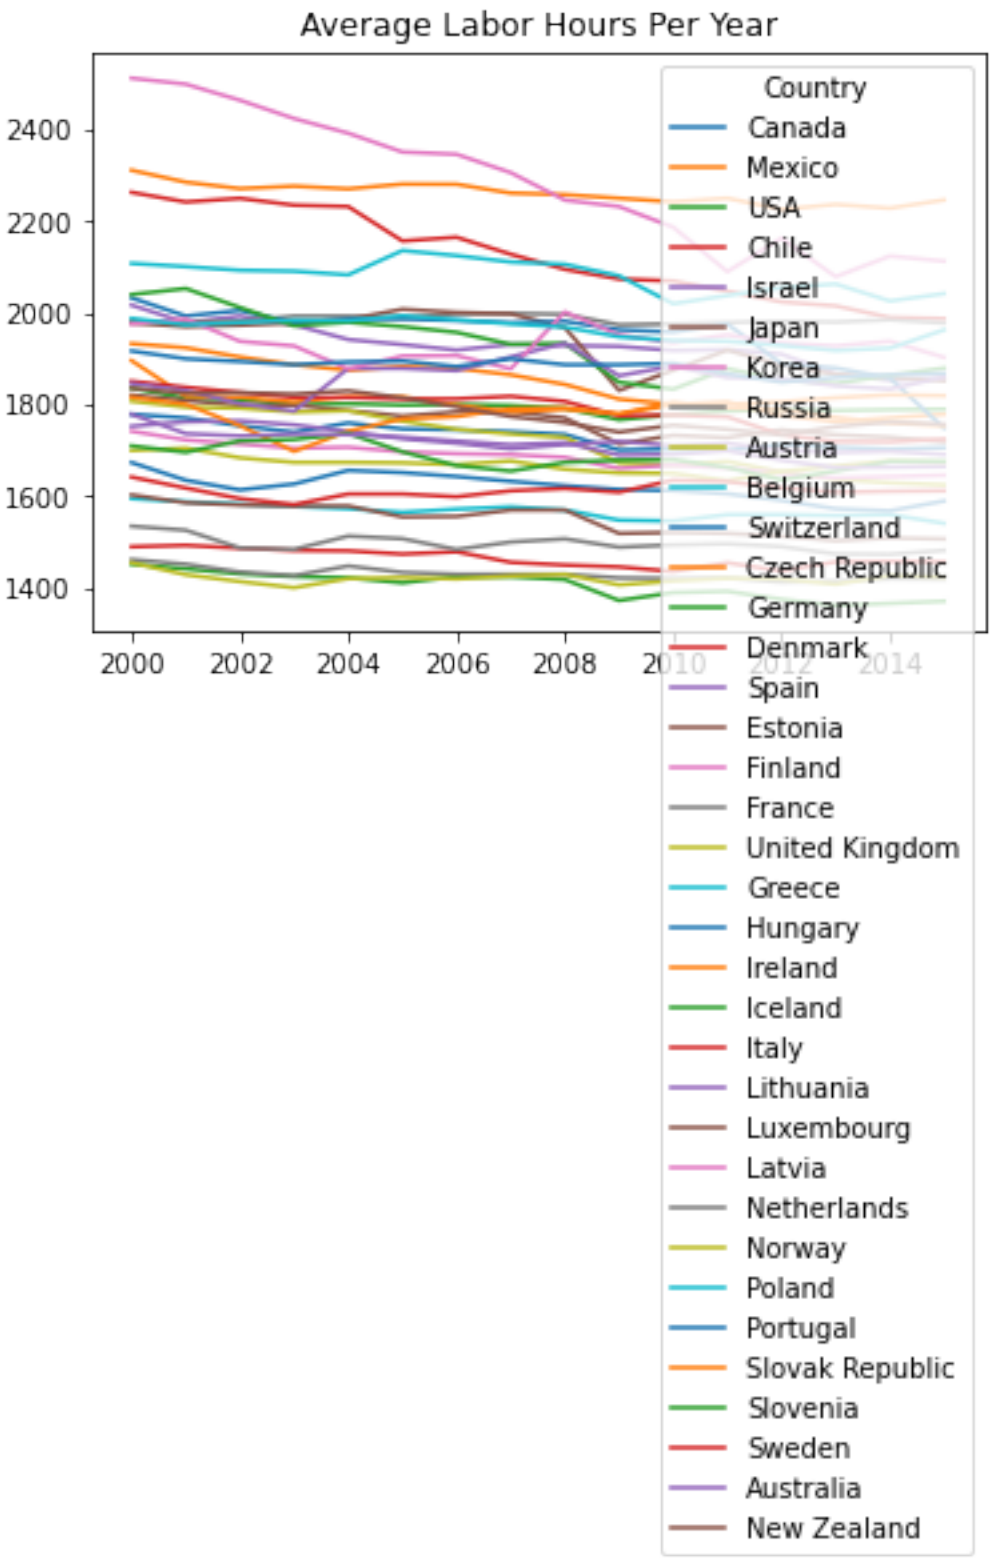
Average Labor Hours Per Year
2400
2200
2000
1800
1600
1400
2000
2002
2004
2006
2008
2010
2012
2014
Country
Canada
Mexico
USA
Chile
Israel
Japan
Korea
Russia
Austria
Belgium
Switzerland
Czech Republic
Germany
Denmark
Spain
Estonia
Finland
France
United Kingdom
Greece
Hungary
Ireland
Iceland
Italy
Lithuania
Luxembourg
Latvia
Netherlands
Norway
Poland
Portugal
Slovak Republic
Slovenia
Sweden
Australia
New Zealand

0.0.13 let us customize this plot, so that country names appear outside the chart

Update plot() with the following features figsize=(10,10), colormap='rainbow', linewidth=2, loc='right'

[56]:
```
world.transpose().plot(figsize=(10,10), colormap='rainbow', linewidth=2,␣
↪title='Average Labor Hours Per Year')
plt.legend(loc='right', bbox_to_anchor=(1.3, 0.5))
plt.show()
```

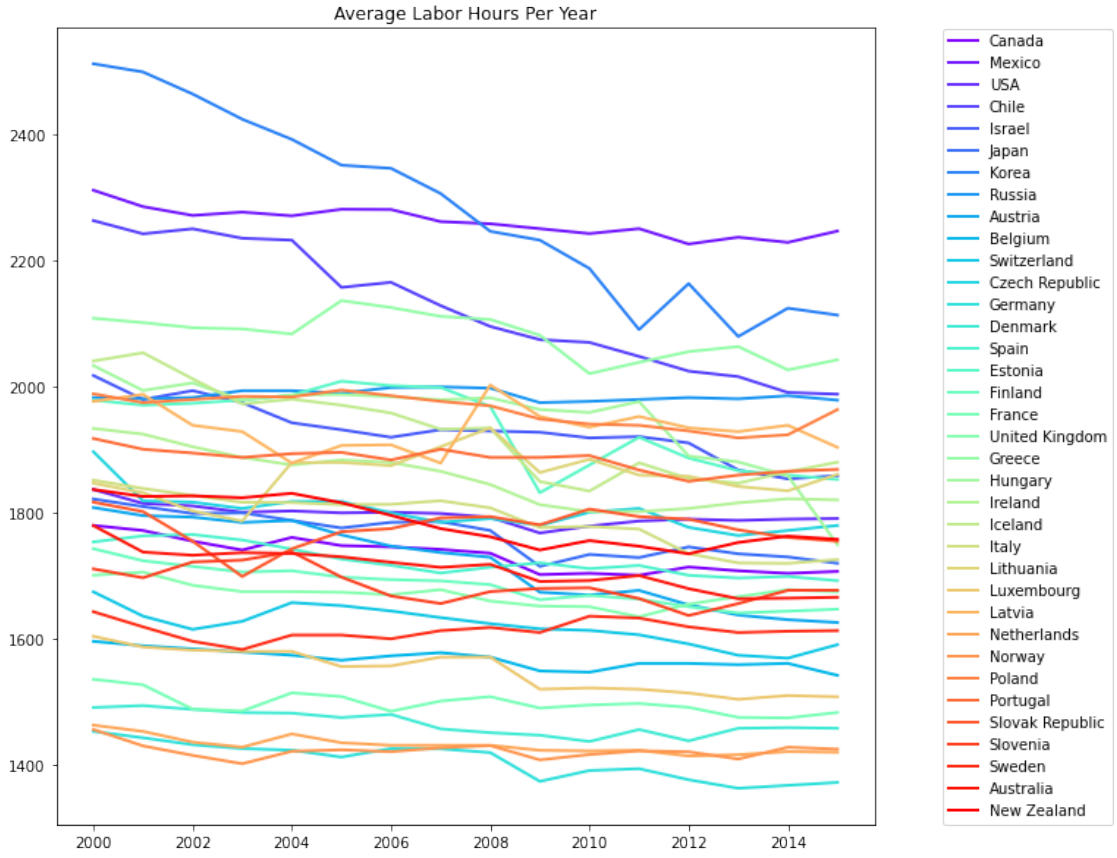

0.0.14 Merging Historical Labor Data

It's nice being able to see how the labor hours have shifted since 2000, but in order to see real trends emerge, we want to be able to see as much historical data as possible. The data collection team was kind enough to send data from 1950 to 2000, let's load it in and take a look.

[57]:
```
historical = pd.read_csv('historical.csv', index_col=0)
historical.head()
```

[57]:
```
             1950  1951  1952  1953  1954  1955  1956  1957  1958  1959  … \
Country                                                                    …
```

```
Australia     NaN   NaN   NaN   NaN   NaN   NaN   NaN   NaN   NaN   NaN  …
Austria       NaN   NaN   NaN   NaN   NaN   NaN   NaN   NaN   NaN   NaN  …
Belgium       NaN   NaN   NaN   NaN   NaN   NaN   NaN   NaN   NaN   NaN  …
Canada        NaN   NaN   NaN   NaN   NaN   NaN   NaN   NaN   NaN   NaN  …
Switzerland   NaN   NaN   NaN   NaN   NaN   NaN   NaN   NaN   NaN   NaN  …

               1990     1991     1992     1993     1994      1995      1996  \
Country
Australia    1779.5  1774.90  1773.70  1786.50  1797.60  1793.400  1782.700
Austria         NaN      NaN      NaN      NaN      NaN  1619.200  1637.150
Belgium      1662.9  1625.79  1602.72  1558.59  1558.59  1515.835  1500.295
Canada       1789.5  1767.50  1766.00  1764.50  1773.00  1771.500  1786.500
Switzerland     NaN  1673.10  1684.80  1685.80  1706.20  1685.500  1658.900

                1997     1998    1999
Country
Australia    1783.600  1768.40  1778.8
Austria      1648.500  1641.65  1654.0
Belgium      1510.315  1513.33  1514.5
Canada       1782.500  1778.50  1778.5
Switzerland  1648.600  1656.60  1678.4

[5 rows x 50 columns]
```

[58]:
```
print("World rows & columns: ", world.shape)
print("Historical rows & columns: ", historical.shape)
```

```
World rows & columns:  (36, 16)
Historical rows & columns:  (39, 50)
```

0.0.15 Merge historical dataframe with world dataframe and store in a new variable, world_historical

[59]:
```
world_historical = pd.merge(historical, world, left_index=True,␣
↪right_index=True, how='right')
```

0.0.16 Print size of world_historical dataframe

[60]:
```
print(world_historical.shape)
```

```
(36, 66)
```

0.0.17 Print top-5 of world_historical dataframe

[61]:
```
world_historical.head()
```

[61]:
```
              1950    1951    1952    1953    1954    1955    1956    1957  \
Country
```

```
Canada      NaN     NaN     NaN     NaN     NaN     NaN     NaN     NaN
Mexico      NaN     NaN     NaN     NaN     NaN     NaN     NaN     NaN
USA      1960.0  1975.5  1978.0  1980.0  1970.5  1992.5  1990.0  1962.0
Chile       NaN     NaN     NaN     NaN     NaN     NaN     NaN     NaN
Israel      NaN     NaN     NaN     NaN     NaN     NaN     NaN     NaN

            1958    1959  …    2006    2007    2008    2009    2010    2011  \
Country                   …
Canada      NaN     NaN   …  1745.0  1741.0  1735.0  1701.0  1703.0  1700.0
Mexico      NaN     NaN   …  2280.6  2261.4  2258.0  2250.2  2242.4  2250.2
USA      1936.5  1947.0   …  1800.0  1798.0  1792.0  1767.0  1778.0  1786.0
Chile       NaN     NaN   …  2165.0  2128.0  2095.0  2074.0  2069.6  2047.4
Israel      NaN     NaN   …  1919.0  1931.0  1929.0  1927.0  1918.0  1920.0

            2012    2013    2014    2015
Country
Canada    1713.0  1707.0  1703.0  1706.0
Mexico    2225.8  2236.6  2228.4  2246.4
USA       1789.0  1787.0  1789.0  1790.0
Chile     2024.0  2015.3  1990.1  1987.5
Israel    1910.0  1867.0  1853.0  1858.0

[5 rows x 66 columns]
```

0.0.18 Joining Historical Data

Now that we've done it the hard way and understand table merging conceptually, let's try a more elegant technique. Pandas has a clean method to join on indexes which is perfect for our situation. ### Use join method to join historical dataframe and world dataframe and store result in world_historical dataframe

[62]:
```
world_historical = historical.join(world, how='right')
world_historical.head()      # Print head of world_historical dataframe
```

[62]:
```
            1950    1951    1952    1953    1954    1955    1956    1957  \
Country
Canada       NaN     NaN     NaN     NaN     NaN     NaN     NaN     NaN
Mexico       NaN     NaN     NaN     NaN     NaN     NaN     NaN     NaN
USA       1960.0  1975.5  1978.0  1980.0  1970.5  1992.5  1990.0  1962.0
Chile        NaN     NaN     NaN     NaN     NaN     NaN     NaN     NaN
Israel       NaN     NaN     NaN     NaN     NaN     NaN     NaN     NaN

            1958    1959  …    2006    2007    2008    2009    2010    2011  \
Country                   …
Canada       NaN     NaN  …  1745.0  1741.0  1735.0  1701.0  1703.0  1700.0
Mexico       NaN     NaN  …  2280.6  2261.4  2258.0  2250.2  2242.4  2250.2
USA       1936.5  1947.0  …  1800.0  1798.0  1792.0  1767.0  1778.0  1786.0
Chile        NaN     NaN  …  2165.0  2128.0  2095.0  2074.0  2069.6  2047.4
```

```
Israel      NaN     NaN  …  1919.0  1931.0  1929.0  1927.0  1918.0  1920.0

           2012    2013    2014    2015
Country
Canada   1713.0  1707.0  1703.0  1706.0
Mexico   2225.8  2236.6  2228.4  2246.4
USA      1789.0  1787.0  1789.0  1790.0
Chile    2024.0  2015.3  1990.1  1987.5
Israel   1910.0  1867.0  1853.0  1858.0

[5 rows x 66 columns]
```

0.0.19 Plot, transposed world_historical dataframe

[63]:
```
world_historical.sort_index(inplace=True)
world_historical.transpose().plot(figsize=(15,10), colormap='rainbow',␣
 ↪linewidth=2, title='Average Labor Hours Per Year')
plt.legend(loc='right', bbox_to_anchor=(1.15, 0.5))
plt.show()
```

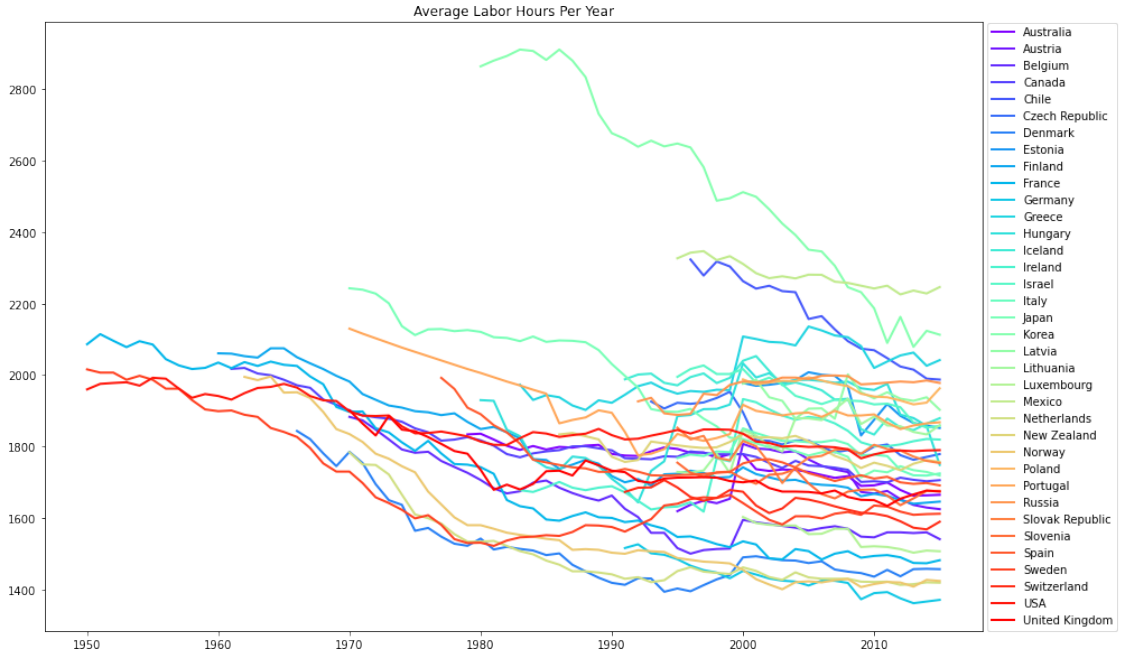

0.0.20 Which country worked longer hours per year?

[65]:
```
work=world.mean(axis=1)

long=max(world.mean(axis=1))
```

```
short=min(world.mean(axis=1))
```

[66]:
```
print("country worked longer hours per year : ",work[work == long].index[0])
```

```
country worked longer hours per year :  Korea
```

0.0.21 Which country worked shorter hours per year?

[67]:
```
print("country worked longer hours per year : ",work[work==short].index[0])
```

```
country worked longer hours per year :  Germany
```